

AMIR SALEEM

8 years | Data Analysis • ML Models • Data Engineering • LLM

03323388625 | amir.saleem.1410h@gmail.com | [linkedin.com/in/amirsaleem1990](https://www.linkedin.com/in/amirsaleem1990) | [amirsaleem1990.github.io](https://github.com/amirsaleem1990)
Islamabad, Pakistan

PROFESSIONAL SUMMARY

Data Scientist / ML Engineer experienced in building scalable ETL pipelines and production ML systems, specializing in cloud environments and large-scale data processing. Delivered impactful solutions that enhanced model performance and streamlined decision-making processes through statistical analysis and machine learning techniques. Focused on collaboration and adaptability to drive data-driven insights and foster innovation in fast-paced settings.

TECHNICAL SKILLS

Programming, Tools & Automation::

Python • SQL • Bash • Selenium • Jupyter • Colab • Cursor • VSCode • DBeaver • Workbench • Jira

Databases & Data Storage::

Database Design • Data Lake • Data Warehouse • Vector Databases • PostgreSQL • MySQL • MongoDB • Redis • ChromaDB • Amazon S3 • Amazon Athena • Amazon Redshift • BigQuery

Data Engineering::

ETL • Reverse ETL • Apache Airflow • Apache Kafka • Feature Engineering • Batch Processing • PySpark • Databricks

Cloud, DevOps & Infrastructure::

AWS • AWS Lambda • AWS CloudWatch • Azure • Azure Functions • Docker • CI/CD • GitHub Actions • Git • Linux • MLflow • Model Deployment • Model Monitoring • Model Versioning

Machine Learning & Statistics::

Scikit-learn • PyTorch • TensorFlow • Keras • XGBoost • Gradient Boosting • PCA • Probability • Neural Networks • NLP • Model Evaluation • MLOps

Generative AI & Agentic Systems::

Large Language Models (LLMs) • Prompt Engineering • RAG • LangChain • Agentic AI • AI Agents • Model Context Protocol (MCP) • Embeddings

Analytics & Visualization::

Pandas • Tableau • R-Shiny • Matplotlib • Seaborn • Excel • Google Sheets • Pivot Tables • Data Mining • Statistical Data Analysis • A/B Testing

Web & API Development::

Web Scraping • Django • FastAPI • GraphQL • SQLAlchemy • Pydantic • Alembic • Celery • Pytest • Postman

WORK HISTORY

Python Data Scientist @ Turing

Remote (USA) | Nov 2023 - Current

- LLM Assessment & Analytics:** **Evaluated** and analyzed **LLM** coding performance by reviewing model responses, categorizing outputs across quality rubrics, analyzing evaluation data, and delivering client-ready reports and annotated benchmark datasets.
- Enterprise Sales & Inventory Data Platform:** Designed, built, and maintained production **ETL pipeline** using **AWS Glue Workflows, Python, and SQL** to ingest and transform data from ERP, CRM, POS, supplier files, and e-commerce systems into an **Amazon Redshift** warehouse; implemented incremental loading and data quality checks to reliably process **~2M records per day** for sales, inventory, customer, and supplier analytics.
- Feature Engineering for Classification:** Improved product category classification accuracy from 68% to 83% (+22% relative) on a balanced 8-class marketplace dataset by engineering description TF-IDF, brand-aggregate, price-bucket, and domain keyword features — validated via ablation testing and 5-fold stratified cross-validation.
- LLM Prompt Engineering & Evaluation:** Enhanced, evaluated, and reviewed a large language model (LLM) support bot using prompt engineering — clearer category rules, few-shot examples, and a held-out test set — reducing LLM error rate by 35%.
- Support Chatbot with RAG & Human Handoff:** Built a **RAG support chatbot** (Python, GPT-4) that auto-resolved 40% of FAQ tickets and drafted replies for agents, cutting support manual effort by 40% while escalating billing disputes, VIP, and low-confidence cases to humans.

Data Engineer (Freelance) @ Medusa

Remote (USA) | Jun 2023 - Feb 2024

- Managed **ETL** processes on **AWS** using **Redshift, Athena, and MySQL** within containerized **AWS Lambda functions** running on **Docker**.
- Integrated with existing **CI/CD** pipelines via **Bitbucket** to enable automated, reliable code deployment workflows.
- Deployed **Docker** images to **AWS** infrastructure and monitored system performance and health using **AWS CloudWatch** dashboards and alerts.
- Implemented **reverse ETL pipelines** for Monday.com using **GraphQL API** integration, enabling seamless bi-directional data synchronization.

Data Scientist @ LoveForData

Karachi, Pakistan | Jan 2023 - Nov 2023

- Designed and implemented end-to-end **ETL workflows orchestrated with Apache Airflow**, integrating **Apache Kafka** for real-time data ingestion into a centralized data warehouse.
- Performed exploratory data analysis (**EDA**) using **Python and SQL**, identifying key business patterns and trends that guided strategic decisions.
- Built and deployed predictive **machine learning models using scikit-learn and PyTorch**, improved recall from 55% to 73% while maintaining constant precision for a default prediction model

Data Analyst @ LoveForData

Karachi, Pakistan | Jan 2021 - Dec 2022

- Developed **Django**-based web applications for interactive data visualization and self-service analytics, increasing stakeholder data accessibility by 50%
- Built and deployed a Python-based **ETL pipeline on Azure Functions**, enabling **serverless**, scalable data processing on the **cloud**.
- Applied **market basket analysis** (association rules) to a restaurant client's transaction data, generating actionable insights that improved menu design, marketing strategy, and customer experience.
- Developed multiple production-grade **web scraping** solutions in Python to automate data collection from external websites.
- Shipped a forecasting model, supporting data-driven planning and operational decisions.

Junior Data Analyst @ LoveForData

Karachi, Pakistan | Apr 2019 - Dec 2020

- Performed comprehensive **data audits** to identify and correct data quality issues, reducing data entry errors by 25% across 5 client databases.
- Built an interactive, real-time dashboard using **R-Shiny** for data visualization and exploratory analysis, enabling stakeholder self-service reporting.
- Developed data workflows to extract, transform, and load data from multiple heterogeneous sources into a structured data warehouse.

Data Science Trainee @ LoveForData

Karachi, Pakistan | Oct 2018 - Mar 2019

- Automated data extraction from 20+ external websites daily.
- Developed an **NLP** pipeline supporting sentiment analysis, text summarization, and question answering using Python-based NLP libraries.

Data Science Intern @ LoveForData

Karachi, Pakistan | Aug 2018 - Sep 2018

- Collected, cleaned, and structured datasets.

CERTIFICATIONS & LANGUAGES

- Certifications:** 50+ certifications in Python, Bash, R, SQL, Excel, Tableau, Machine Learning, Data Engineering, Linux, and Statistics. (Visit the [portfolio](#) for full list).
- Languages:** English (Fluent), Urdu (Native), Arabic (Conversational).