

AMIR SALEEM

8 years | Data Analysis • ML Models • Data Engineering • LLM

DATA SCIENTIST | ML ENGINEER | DATA ENGINEER

+92 332 3388625 | amirsaleem1990@hotmail.com | linkedin.com/in/amirsaleem1990 | amirsaleem1990.github.io
Islamabad, Pakistan

PROFESSIONAL SUMMARY

Data Scientist / ML Engineer with 8 years of experience building scalable ETL pipelines and production ML systems. Specialized in LLM fine-tuning, cloud (AWS/GCP), and large-scale data processing. Delivered solutions improving model performance, reducing manual effort, and enabling data-driven decision making.

TECHNICAL SKILLS

Programming, Tools & Automation:: Python • SQL • Bash • R • Selenium • Jupyter • Colab • Cursor • VSCode • DBeaver • Workbench • Jira	Databases & Data Storage:: Database Design • Data Lake • Data Warehouse • Vector Databases • PostgreSQL • MySQL • MongoDB • Redis • ChromaDB • Amazon S3 • Amazon Athena • Amazon Redshift • BigQuery
Data Engineering:: ETL • ELT • Reverse ETL • Feature Engineering • Batch Processing • PySpark • Databricks	Cloud, DevOps & Infrastructure:: AWS • AWS Lambda • AWS CloudWatch • Azure • Azure Functions • Docker • CI/CD • GitHub Actions • Git • Linux • MLflow • Model Deployment • Model Monitoring • Model Versioning
Machine Learning & Statistics:: Scikit-learn • PyTorch • TensorFlow • Keras • XGBoost • Gradient Boosting • PCA • Probability • Neural Networks • NLP • Model Evaluation • MLOps	Generative AI & Agentic Systems:: Large Language Models (LLMs) • Prompt Engineering • RAG • LangChain • Agentic AI • AI Agents • Model Context Protocol (MCP) • Embeddings
Analytics & Visualization:: Pandas • Tableau • R-Shiny • Matplotlib • Seaborn • Excel • Google Sheets • Pivot Tables • Data Mining • Statistical Data Analysis • A/B Testing	Web & API Development:: Web Scraping • Django • FastAPI • GraphQL • SQLAlchemy • Pydantic • Alembic • Celery • Pytest • Postman

PROFESSIONAL EXPERIENCE (8 YEARS)

Python Data Scientist @ Turing - Improved model accuracy by 22% using advanced feature engineering. - Enhanced, evaluated, and reviewed large language model (LLM) through prompt engineering techniques, reducing LLM error rate by 35% - Designed and maintained high-throughput ETL pipelines and data workflows using Python and SQL, processing ~ 5M+ records daily with improved reliability. - Eliminated repetitive data processing tasks using Python scripting, which reduced manual effort by 40%. - Partnered with cross-functional teams (engineering, product, and business) to convert data insights into actionable product and strategy improvements. - Optimized existing data pipelines for scalability and performance, enabling faster iteration cycles for the data science team.	<i>Remote (USA) Nov 2023 - Mar 2026</i>
Data Engineer (Freelance) @ Medusa - Built and managed ETL processes on AWS using Redshift, Athena, and MySQL within containerized AWS Lambda functions running on Docker. - Integrated with existing CI/CD pipelines via Bitbucket to enable automated, reliable code deployment workflows. - Deployed Docker images to AWS infrastructure and monitored system performance and health using AWS CloudWatch dashboards and alerts. - Implemented reverse ETL pipelines for Monday.com using GraphQL API integration, enabling seamless bi-directional data synchronization.	<i>Remote (USA) Jun 2023 - Feb 2024</i>
Data Scientist @ LoveForData - Designed and implemented End-to-End ETL workflows to extract, transform, and load data from diverse sources into a centralized data warehouse. - Performed exploratory data analysis (EDA) using Python and SQL, identifying key business patterns and trends that guided strategic decisions. - Built and deployed predictive machine learning models using scikit-learn and PyTorch, improved recall from 55% to 73% while maintaining constant precision for a default prediction model	<i>Karachi, Pakistan Jan 2023 - Nov 2023</i>
Data Analyst @ LoveForData - Developed Django-based web applications for interactive data visualization and self-service analytics, increasing stakeholder data accessibility by 50% - Built and deployed a Python-based ETL pipeline on Azure Functions, enabling serverless, scalable data processing on the cloud. - Applied market basket analysis (association rules) to a restaurant client's transaction data, generating actionable insights that improved menu design, marketing strategy, and customer experience. - Developed multiple production-grade web scraping solutions in Python to automate data collection from external websites. - Shipped a forecasting model, supporting data-driven planning and operational decisions.	<i>Karachi, Pakistan Jan 2021 - Dec 2022</i>
Junior Data Analyst @ LoveForData - Performed comprehensive data audits to identify and correct data quality issues, reducing data entry errors by 25% across 5 client databases. - Built an interactive, real-time dashboard using R-Shiny for data visualization and exploratory analysis, enabling stakeholder self-service reporting. - Developed data workflows to extract, transform, and load data from multiple heterogeneous sources into a structured data warehouse.	<i>Karachi, Pakistan Apr 2019 - Dec 2020</i>
Data Science Trainee @ LoveForData - Automated data extraction from 20+ external websites daily. - Developed an NLP pipeline supporting sentiment analysis, text summarization, and question answering using Python-based NLP libraries.	<i>Karachi, Pakistan Oct 2018 - Mar 2019</i>
Data Science Intern @ LoveForData Collected, cleaned, and structured datasets.	<i>Karachi, Pakistan Aug 2018 - Sep 2018</i>

CERTIFICATIONS & LANGUAGES

- Certifications:** 50+ certifications in Python, Bash, R, SQL, Excel, Tableau, Machine Learning, Data Engineering, Linux, and Statistics. (Visit the [portfolio](#) for full list).
- Languages:** English (Fluent), Urdu (Native), Arabic (Conversational).